

Privacy-Preserving Multi-hop Payments with Constant Collateral

1st Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
 City, Country
 email address or ORCID

2nd Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
 City, Country
 email address or ORCID

Abstract—This document is a model and instructions for L^AT_EX.

Index Terms—component, formatting, style, styling, insert

I. INTRODUCTION

In recent years, permissionless cryptocurrencies, have emerged as a novel means to facilitate secure and reliable payments within a decentralized framework, garnering significant attention from both academia and industry. These cryptocurrencies employ a consensus mechanism to verify each transaction, which is then recorded on a publicly distributed ledger known as blockchain. Unfortunately, the widespread adoption of cryptocurrencies is hindered by notable scalability challenges. Complex consensus mechanisms, like Bitcoin’s Proof-of-work(PoW), and the limited block size of the blockchain contribute to the issue. The theoretical throughput of Bitcoin stands at approximately 10 transactions per second(TPS), with a transaction confirmation time of around 1 hour. In contrast, traditional decentralized payment networks, such as Visa, boast the capability up to 47,000 TPS. Furthermore, the presence of high transaction fees renders small-value payments impractical for cryptocurrency users.

One promising solution proposed to address the scalability issue is the utilization of payments channels(PCs), which facilitate quick and validated transactions between two parties off-chain. In a nutshell, a payment channel is represented as a multi-signature address on the blockchain, with both parties initially depositing a certain amount of funds. Subsequently, during each transaction, the parties adjust the fund allocation by generating and exchanging signed transaction messages off-chain. When the participants decide to settle the channel or encounter a dispute, they initiate the closing process by broadcasting the latest signed transaction to the blockchain.

However, this straightforward construction forces users to establish a separate payment channel with each potential transaction partner and lock coins for each one. Additionally, these coins cannot be used for other things until the channel is closed. Payment channel network(PCN) mitigates this issue by enabling transactions between two users who do not share a direct payment channel. This is achieved through multi-hop payments(MHP) by leveraging intermediaries and payment channels within the PCN.

A. PCN based on HTLC and some limitations

Currently, the most renowned and valuable PCN system is the Lightning Network(LN), which is implemented on the Bitcoin. To guarantee the security and atomicity of multi-hop payments, LN employs a specialized mechanism called hash timed-locked contracts(HTLCs). The HTLC mechanism involves two phases. In the locking phase, coins are locked in each channel using a hash chosen by the receiver from the left to right. In the subsequent release phase, the receiver present the preimage of the hash value he initially selected to unlock coins and all the intermediaries can claim coins by providing the same preimage. An example of a MHP based HTLC is shown in figure 1. Unfortunately, PCNs based on HTLC have the following fundamental problems.

Collateral To ensure the atomic execution of MHP, each HTLC needs to specify an expiration time(also called collateral time) for the locked coins. Once time expires, the locked coins will be returned to the user on the left side of the channel. Moreover, the collateral time increases gradually from left to right in a staggered fashion. Assuming α coins are being paid through n PCs, the collateral time for PC i is $t_i \geq t_i + \Delta$, $i = 1, 2, \dots, n$, Δ denotes the maximum duration between when a transaction is broadcast and when it is confirmed transaction by the blockchain. This staggered manner ensures that each user has enough time to await the timely update of his right channel and expeditiously broadcast the state of his left channel to the blockchain, regardless of whether the counterparty cooperates.

While this mechanism effectively prevent honest users losing coins, it is essential to note that the overall collateral for the entire payment path is $\Theta(n^2 \cdot \alpha \cdot \Delta)$ in the worst-case




Fig. 1. An MHP from Alice to Bob for value α coins using HTLC contracts. A HTLC contract denotes as HTLC(X, Y, h, α , t) means:(i)If Y can present a value r such that $H(r) = h$ before timeout t , X pays α coins to Y.(ii)If timeout t expires, X gets back the α locked coins

scenario in units $\text{coins} \times \text{time}$. In the event of a failure of the MHP where all coins are returned, users cannot receive any benefit but incur potential financial losses. This is due to the fact that the locked coins cannot be utilized for other transactions during this period. Furthermore, the high volatility of cryptocurrency prices can exacerbate these economic losses, further magnifying the impact on users' financial well-being. Simultaneously, it is crucial to acknowledge that this issue is not limited to the HTLC mechanism alone, but rather applies to all multi-hop payments protocols within two-phase communication [7].

Anonymity Despite the employment of onion routing [8], [9] for anonymity in the LN, it is important to be aware that the information contained in HTLCs can inadvertently expose details about payment path. In our analysis, we examine the inherent limitations of the HTLC mechanism regarding anonymity from both *off-chain* and *on-chain* perspectives. These aspects are critical consideration for the development of proposed PCNs.

II. BACKGROUND

In this section, we provide an overview on the background and the notations used throughout the paper.

A. UTXO model

In this work, we assume the underlying blockchain, like Bitcoin, is based on the UTXO model. Transaction output is the fundamental component of Bitcoin transaction, which is an indivisible Bitcoin currency recorded on the blockchain and recognized as valid by the entire network. The Bitcoin complete nodes tracks all available outputs called *unspent transaction output* (UTXO). UTXOs can be any value and once generated are indivisible, a UTXO can only be consumed as a whole in a single transaction. The output consists of two parts, which we represent as a tuple $\theta := \{\text{cash}, \phi\}$, $\theta.\text{cash}$ is the output value, and the $\theta.\phi$ is the condition to spend this output, it also called *locking script*. $\text{Onesig}(U)$ indicates that the condition required to spend this output is a digital signature. We say that a user U can spend an output only if $\theta.\phi$ contains only a signature w.r.t verification U 's public key, if multiple signatures are required, we use $\text{MultiSig}(U_1, U_2, \dots, U_n)$.

Transactions consume previously recorded unused UTXOs and create new UTXOs available for future transaction, in this way, the transaction continues in the form of a chain of owners. on the chain, the input of the transaction corresponds to the output of the previous transaction. We denote a transaction as a tuple $\text{TX} := \{\text{id}, \text{input}, \text{output}, \text{timelock}, \text{witness}\}$, $\text{TX.id} \in \{0, 1\}^*$ is the identifier of a transaction and $\text{TX.id} = \mathcal{H}(\text{TX.input}, \text{TX.output}, \text{TX.timelock})$, \mathcal{H} is a hash function, modeled as a random oracle. TX.input and TX.output denotes the list of the inputs and the list of new outputs respectively. TX.timelock defined as the earliest time a transaction is valid and can be transmitted on the network or added to the blockchain, it defaults to 0 in most transactions. $\text{TX.witness} \in \{0, 1\}^*$, also called *ScriptSig*, is part of the transaction input to address or satisfy the spending conditions set by the *locking*




Fig. 2. Left transaction TX_1 is recorded on the blockchain, which has two output. Value α_1 can be spent by A, with a transaction signed w.r.t PK_A after t_1 rounds, and the output of value α_2 can be spent by a transaction signed w.r.t PK_A and PK_B but only if at least t_2 rounds has passed through after the transaction being published on the blockchain. The right transaction TX_2 has one input which is the second output of TX_1 with the value α_1 . TX_2 has only one output containing α_3 coins, which can spend by a transaction whose witness satisfies the condition $\varphi_1 \vee (\varphi_2 \wedge \varphi_3)$

script on the output. Actually, before being recorded on the blockchain, transactions must go through the consensus mechanism of all nodes. During this period, each node will independently verify the transaction. Transactions need to be verified as follows: For a transaction(1) The sum of the input value cannot be less than the sum of the output value;(2)For each input, the quoted output cannot exist in any other transaction, it must exist and not be spent;(3) The *ScriptSig* for each input must be validated against the *locking script* for the corresponding output. To put it simply, the transaction must provide valid validation that satisfies the spending conditions of each input.

We use charts to visualize the transaction for a clearer illustration. Rounded rectangles represent transactions, thick-edged rectangles represent transactions that have already been published on the blockchain, and thin-edged rectangles for transactions to be published. The transaction contains at least one box to represent the output of the transaction, and the value in the box indicates the number of coins in this output. On the arrows coming from the output, are noted the conditions under which this output is spent. The public key below the arrow indicates who can use this output; Above the arrow is the timelock for the output(in this work only uses timelock as additional condition, which in practice could be any script supported by the underlying blockchain scripting language). There are two types of timelocks: *relative time lock* and *absolute time lock*, we use $+t$ to represent the relative time lock, that is, the transaction only valid at least t blocks has passed through after the transaction being recorded on the blockchain; The absolute time lock $\geq t$ specifies the absolute time point, indicating that the first transaction has passed through t round after being recorded on the blockchain. Finally, we use a diamond to represent the relationship of "or" that the output conditions are different, expressed in the symbolic form as $\varphi = \varphi_1 \vee \dots \vee \varphi_n$, where φ is the output locking script, written on the arrows of the output. A complete example is given in Fig. 1.

B. Payments channels

The payment channel is opened by two users locking some coins on the blockchain, and then the two parties can make as many quick instant confirmation transactions off-chain, that is, without waiting for transmitting the transaction to the blockchain, as long as the total amount of each transaction is less than the locked value. During the lifetime of the channel, only two transactions will be recorded on the blockchain, one is funding transaction TX_f which used to open the channel, the other is settlement transaction for close the channel. The funding transaction TX_f determines the balance of this channel, and the channel is open when TX_f is broadcast to the network and recorded on the blockchain, both parties update the channel balance by exchanging signed transactions off-chain, these transactions holding the reassigned balances are called state of the channel. Finally, the channel is closed by posting settlement transaction, which is the final state, on the blockchain. In a more detail, there are three operations of the payment channel: *open*, *update*, and *close*.

1) *Open*: Suppose that Alice and Bob want to create a payment channel, and to do so, Alice and Bob agree on a funding transaction TX_f . The output value of TX_f is $\alpha_A + \alpha_B$, which is the total amount of the channel, and there are two inputs to the TX_f are from Alice and Bob respectively, α_A represents the initial deposit provided by Alice in the channel, and α_B represents Bob's. When TX_f is published on the blockchain, it indicates that the payment channel has been opened.

2) *update*: After the channel is opened, to update the balance status of both sides of the channel, that is to make a transaction, parties exchange signatures on the commit transaction. Each time a new "commit transaction" is exchanged, the previous state of the channel is overwritten and updated to the newer state, while the old state is saved locally. Therefore, only the most recent commit transaction can be "executed". To ensure the security of payments, payment channel use a penalty mechanism to: 1) prevent malicious user from closing the channel by sending an outdated but self-advantageous state to the chain, and 2) enable honest user to acquire all the coins locked in the channel when dishonest user post outdated promised transactions. We use the recent work *generalized channels* [2] as the basic structure, which split the "commit transaction" into two parts, the commit transaction TX_c responsible for penalizing and the split transaction TX_s , that hold the actual output of the channel. assume that Alice wants to pay $x_A \leq \alpha_A$ coins to Bob. To do this, they create new commit transaction TX'_c and split transaction TX'_s , at the same time, they pick the new *secret* R' to represent the commitment of the two parties to the new balance of the channel, then they exchange the *revocation secret* r of the previous round of transactions TX_c and TX_s , after that they make the new commit transaction TX'_c and split transaction TX'_s overwrite previous ones. TX'_s represents the updated balance of the channel, and both parties save the previous round of transactions (TX_c and TX_s) with the *revocation*




Fig. 3. Details of generalized channel, $(x_1, \varphi_1) \dots (x_1, \varphi_1)$ are the states of the channel. The output of TX_c is three different conditions, two of which are punishment mechanisms used to prevent dishonest users. The condition ϕ_B is punishment for B, represents the verification of B' revocation secret, When A discovers that B submit an old transaction, it can take all the money in the channel by B' revocation secret, which B has exchanged for A when the old transaction is updated. In the same way, ϕ_A represents the verification of A' revocation secret that can punish A. The value of Δ is time upper bound that takes to publish a transaction on a blockchain.

secret r of the secret value in local memory.

3) *close*: Finally, assume both parties agree to close the channel, They can collaboratively submit $\overline{TX_c}$ and $\overline{TX_s}$, which represent the last transaction they agreed, to the blockchain to close the channel. If the dishonest one submit an expired promised transaction, the malicious party can be penalized by closing the channel through the revocation mechanism unilaterally.

C. Payment channel networks

Suppose that U_0 now wants to transfer some coins to U_2 , but there is no direct payment channel between U_0 and U_2 , and instead, there is an indirect path in which U_1 participates, that is, U_0 has a open payment channel with U_1 , U_1 has an open payment channel with U_2 , then U_0 can safely transfer coins to U_2 through this path. This way of allowing the atomic transfer of coins from the sender to the receiver through intermediaries atomically called *multi-hop payment* (MHP). A Payment Channel Network (PCN) [11] is a graph consisted of users as the vertices and edges as channels between pairs of users. It can make a large number of single channels connected in series, thus forming an interconnected, vast payment network. Any user on the PCN can pay in MHP through a path of open payment channels.

HTLC However, in the payment process, there is no way to ensure that each user can perform honestly, and malicious users will cause honest users to lose coins. To guarantee the atomicity of payments, the Lightning Network (LN) [3] is done by using a Hash Time-Lock Contracts (HTLC), which works as follows. In the payment channel of u_0 and u_1 , u_0 locks her coins in the output, and this output can be spent on two conditions: 1) beyond the predefined time, u_0 will get her money back; 2) once u_1 provides a pre-image r_A with a hash value $\mathcal{H}(r_A)$ set by u_0 , then u_1 can take the money. In a nutshell, sender u_0 who wants to pay α coins to the receiver u_n through some intermediary $\{u_i\}_{i \in [1, n-1]}$, and two users u_j and u_{j+1} for $j \in [0, n-1]$ have an open payment channel. First, u_n selects a random number r , calculates its hash $y := \mathcal{H}(r)$ and sends y to u_0 . Second, the u_0 creates

a new state with three outputs $(\theta_1, \theta_2, \theta_3)$ and establishes the HTLC with u_1 . The value of θ_1 is α , θ_2 is the balance of u_0 minus α and θ_3 is the balance of u_1 , where θ_1 is the output used by HTLC. HTLC stipulates that once u_1 provides a pre-image of x such that $y = \mathcal{H}(x)$, then u_1 can take the coins, or return them to u_0 if the timeout after $n \cdot T$ has expired. Then u_1 repeats this step for her right neighbor u_2 , again using y but decreasing time, which is $(n-1)T$. Repeat this step until the receiver u_n is reached with a timeout of T . This process is called the *setup phase*.

During the *open phase*, the receiver u_n can present r (the secret value of giving coins to u_n requested in HTLC) to her left neighbor u_{n-1} . After that, both parties agree to update their channel to a new state off-chain, and the balance of u_n will increase α coins, or u_n can publish the state on-chain and a transaction with witness r spending the money from the HTLC to itself. The other intermediate users use r to continue the process, u_i reveals the secret r to its left neighbor u_{i-1} and opens the HTLC. But for this process, enough time needs to be reserved for u_{i-1} to ask her left neighbor for her coins, otherwise, u_i can claim the money of HTLC by spending the HTLC on-chain at the last possible moment, due to blockchain latency, user u_{i-1} will notice that too late and will no longer be able to claim HTLC money from u_{i-2} . This is the reason of time-locks on HTLC are interleaved, i.e. increasing from right to left.

According to the previously mentioned, MHPs done using HTLC requires *setup phase* and *open phase* two rounds of paired sequential communication, and its collateral lock time is linearly with the path length. As mentioned in [4], such a setting may trigger a grieving attack. Secondly, in MHPs, in order to incentivize the intermediary to participate in helping the sender and receiver complete the payment, an additional fee must be given to the intermediary at each payment. A wormhole attack [5] allows two colluding users to steal their fees by skipping honest intermediary in payments.

Blize Blize [6] improves MHPs with a transaction which acts as a global event and a *pay-unless-revoke* paradigm. On the basis of ensuring security when malicious intermediaries exist, it only requires one round of communication through the path and constant collateral lock time. Specifically, sender u_0 creates a unique global transaction *emphEnableRefund*, denoted by TX_{er} . If all channels from u_0 to receiver u_n are successfully updated, u_n will send a confirmation to u_0 , and coins will naturally be paid to u_n through these updated channels. If any channel update fails (e.g. an intermediary offline), u_0 will publish TX_{er} to the blockchain before the preset time T to trigger all refunds. To achieve this, each party u_i for $i \in [0, n-1]$ creates an output of \mathbb{N} that is spendable in two ways: 1) u_{i+1} claim it after time T , or 2) if TX_{er} is on the blockchain before time T , the coins can be returned to u_i . **Thora** The same idea of global transaction is applied in Thora [7], but it proceeds logic is reversed of Blize, which is considered *revoke-unless-pay* paradigm. Thora supports multiple senders and receivers that do not need to be connected to each other, breaking the path-based topology, and com-

plete atomic updates of arbitrary channels. In this Protocol, each receiver creates her own TX_{ep} , the global transaction *Enable Payment*, each TX_{ep} has outputs to all receivers. All receivers send their TX_{ep} to all other parties, and this time each sender creates one TX_p per TX_{ep} . Then, if all channels are updated successfully, the senders can transfer coins to the receivers. If some transfer fails, the receivers can post TX_{ep} on the chain and force all payments. This allows 1) the sender to return her coins if at least one channel fails to execute. All refunds possible only if a timeout T (a preset parameter) expires, so all senders can refund their coins if the coins have not been spent by the receivers after T , and 2) if all payments are successful, the receiver to claim her coins. For each channel, the sender updates the channel and creates a TX_p , which transaction can transfer coins to the receiver only after TX_{ep} appeared on the blockchain before time T . If at least one receiver does not receive the coins(channel fails to execute the payment), TX_{ep} will be triggered, and all payment transactions will be enforced before time T .

III. SOLUTION OVERVIEW

In this section, we present our key idea.

A. Security and privacy goals

B. Key idea

IV. CONSTRUCTION

A. Building blocks

B. Protocol description

V. SECURITY ANALYSIS

VI. EVALUATION

VII. CONCLUSION

Conclude the paper.

REFERENCES

Please number citations consecutively within brackets [1]. The sentence punctuation follows the bracket [2]. Refer simply to the reference number, as in [3]—do not use “Ref. [3]” or “reference [3]” except at the beginning of a sentence: “Reference [3] was the first . . .”

Number footnotes separately in superscripts. Place the actual footnote at the bottom of the column in which it was cited. Do not put footnotes in the abstract or reference list. Use letters for table footnotes.

Unless there are six authors or more give all authors’ names; do not use “et al.”. Papers that have not been published, even if they have been submitted for publication, should be cited as “unpublished” [4]. Papers that have been accepted for publication should be cited as “in press” [5]. Capitalize only the first word in a paper title, except for proper nouns and element symbols.

For papers published in translation journals, please give the English citation first, followed by the original foreign-language citation [6].

REFERENCES

- [1] G. Eason, B. Noble, and I. N. Sneddon, "On certain integrals of Lipschitz-Hankel type involving products of Bessel functions," *Phil. Trans. Roy. Soc. London*, vol. A247, pp. 529–551, April 1955.
- [2] J. Clerk Maxwell, *A Treatise on Electricity and Magnetism*, 3rd ed., vol. 2. Oxford: Clarendon, 1892, pp.68–73.
- [3] I. S. Jacobs and C. P. Bean, "Fine particles, thin films and exchange anisotropy," in *Magnetism*, vol. III, G. T. Rado and H. Suhl, Eds. New York: Academic, 1963, pp. 271–350.
- [4] K. Elissa, "Title of paper if known," unpublished.
- [5] R. Nicole, "Title of paper with only first word capitalized," *J. Name Stand. Abbrev.*, in press.
- [6] Y. Yorozu, M. Hirano, K. Oka, and Y. Tagawa, "Electron spectroscopy studies on magneto-optical media and plastic substrate interface," *IEEE Transl. J. Magn. Japan*, vol. 2, pp. 740–741, August 1987 [Digests 9th Annual Conf. Magnetism Japan, p. 301, 1982].
- [7] L. Aumayr, P. Moreno-Sanchez, A. Kate, and M. Maffei, "Blitz:Secure Multi-Hop Payments Without Two-Phase Commmits," in *USENIX Security*, 2021.
- [8] J. Camenisch and A. Lysyanskaya, "A formal treatment of onion routing," in *CRYPTO*, 2005, pp. 169–187.
- [9] G. Danezis and I. Goldberg, "Sphinx: A Compact and Provably Secure Mix Format," in *IEEE S&P*, 2009.